



American International University- Bangladesh

Department of Computer Science

Faculty of Science & Technology (FST)

Fall: 23-24

DATA WAREHOUSING AND DATA MINING [B]

Project Name: Heart Attack Analysis & Prediction. (Feature selection by ARM)

Name: Md. Minhazul Bari Fahim

ID: 20-42176-1; **Section:** B

Project Overview:

The project involves the analysis and prediction of heart attacks based on a dataset containing various health-related features. The dataset includes information such as age, sex, exercise-induced angina, number of major vessels, chest pain type, resting blood pressure, cholesterol levels, fasting blood sugar, resting electrocardiographic results, maximum heart rate achieved, and a target variable indicating the likelihood of a heart attack (0 for less chance, 1 for more chance). The goal is to build a predictive model that can accurately classify individuals into these two categories, helping in early detection and prevention of heart attacks. The project will utilize Association Rule Mining (ARM) for feature selection, aiming to identify the most relevant features that contribute to the prediction of heart attacks. The selected features will then be used to train machine learning models for heart attack prediction.

Dataset Overview:

The dataset consists of health-related attributes collected from individuals, with a focus on factors associated with heart health. The features include demographic information such as age and sex, lifestyle indicators like exercise-induced angina and fasting blood sugar, as well as medical measurements like resting blood pressure, cholesterol levels, and maximum heart rate achieved. The dataset is labelled, with the target variable indicating whether an individual has a higher or lower chance of experiencing a heart attack. The diverse set of features provides a comprehensive overview of potential risk factors associated with heart health. The use of Association Rule Mining (ARM) will help in identifying patterns and relationships among these features, contributing to the selection of key variables for building an effective predictive model for heart attack risk assessment.

Source: <https://www.kaggle.com/datasets/rashikrahmanpritom/heart-attack-analysis-prediction-dataset/data>

Importing all the necessary libraries and then loading our dataset into jupyter notebook.

```
✓ 1s [1] import numpy as np
      import pandas as pd
      import os
      for dirname, _, filenames in os.walk('/kaggle/input'):
          for filename in filenames:
              print(os.path.join(dirname, filename))
```

```
✓ [74] pip install scikeras
```

```
✓ [73] pip install kmodes
```


```
[62] import matplotlib.pyplot as plt
      import seaborn as sns
      from mlxtend.frequent_patterns import apriori
      from mlxtend.frequent_patterns import association_rules
      from mlxtend.preprocessing import TransactionEncoder
      from sklearn.model_selection import train_test_split
      import time
      from sklearn.metrics import accuracy_score
```


```
# import libraries of models
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from scikeras.wrappers import KerasClassifier
from sklearn.linear_model import LogisticRegression
from kmodes.kmodes import KModes
```

```
# keras
from keras.models import Sequential
from keras.layers import Dense
```


```
[70] df_const = pd.read_csv('heart.csv')
      df = df_const.copy(deep=True)
```


+ Overview of the dataset

✓ 0s  `df_const.info()`

 `<class 'pandas.core.frame.DataFrame'>`
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
Column Non-Null Count Dtype
--- -
0 age 303 non-null int64
1 sex 303 non-null int64
2 cp 303 non-null int64
3 trtbps 303 non-null int64
4 chol 303 non-null int64
5 fbs 303 non-null int64
6 restecg 303 non-null int64
7 thalachh 303 non-null int64
8 exng 303 non-null int64
9 oldpeak 303 non-null float64
10 slp 303 non-null int64
11 caa 303 non-null int64
12 thall 303 non-null int64
13 output 303 non-null int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

+ Missing Values

 `df_const.isnull().sum()`

 `/usr/local/lib/python3.10/dist-packages/ipykernel/i`
`and should_run_async(code)`
age 0
sex 0
cp 0
trtbps 0
chol 0
fbs 0
restecg 0
thalachh 0
exng 0
oldpeak 0
slp 0
caa 0
thall 0
output 0
dtype: int64

The data is very clean already and there are no missing data.

Sample







```
df_const.sample(10)
```


```
/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async`  
and `should_run_async(code)`
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
14	58	0	3	150	283	1	0	162	0	1.0	2	0	2	1
141	43	1	0	115	303	0	1	181	0	1.2	1	0	2	1
121	59	1	0	138	271	0	0	182	0	0.0	2	0	2	1
19	69	0	3	140	239	0	1	151	0	1.8	2	2	2	1
45	52	1	1	120	325	0	1	172	0	0.2	2	0	2	1
154	39	0	2	138	220	0	1	152	0	0.0	1	0	2	1
251	43	1	0	132	247	1	0	143	1	0.1	1	4	3	0
89	58	0	0	100	248	0	0	122	0	1.0	1	0	2	1
163	38	1	2	138	175	0	1	173	0	0.0	2	4	2	1
94	45	0	1	112	160	0	1	138	0	0.0	1	0	2	1

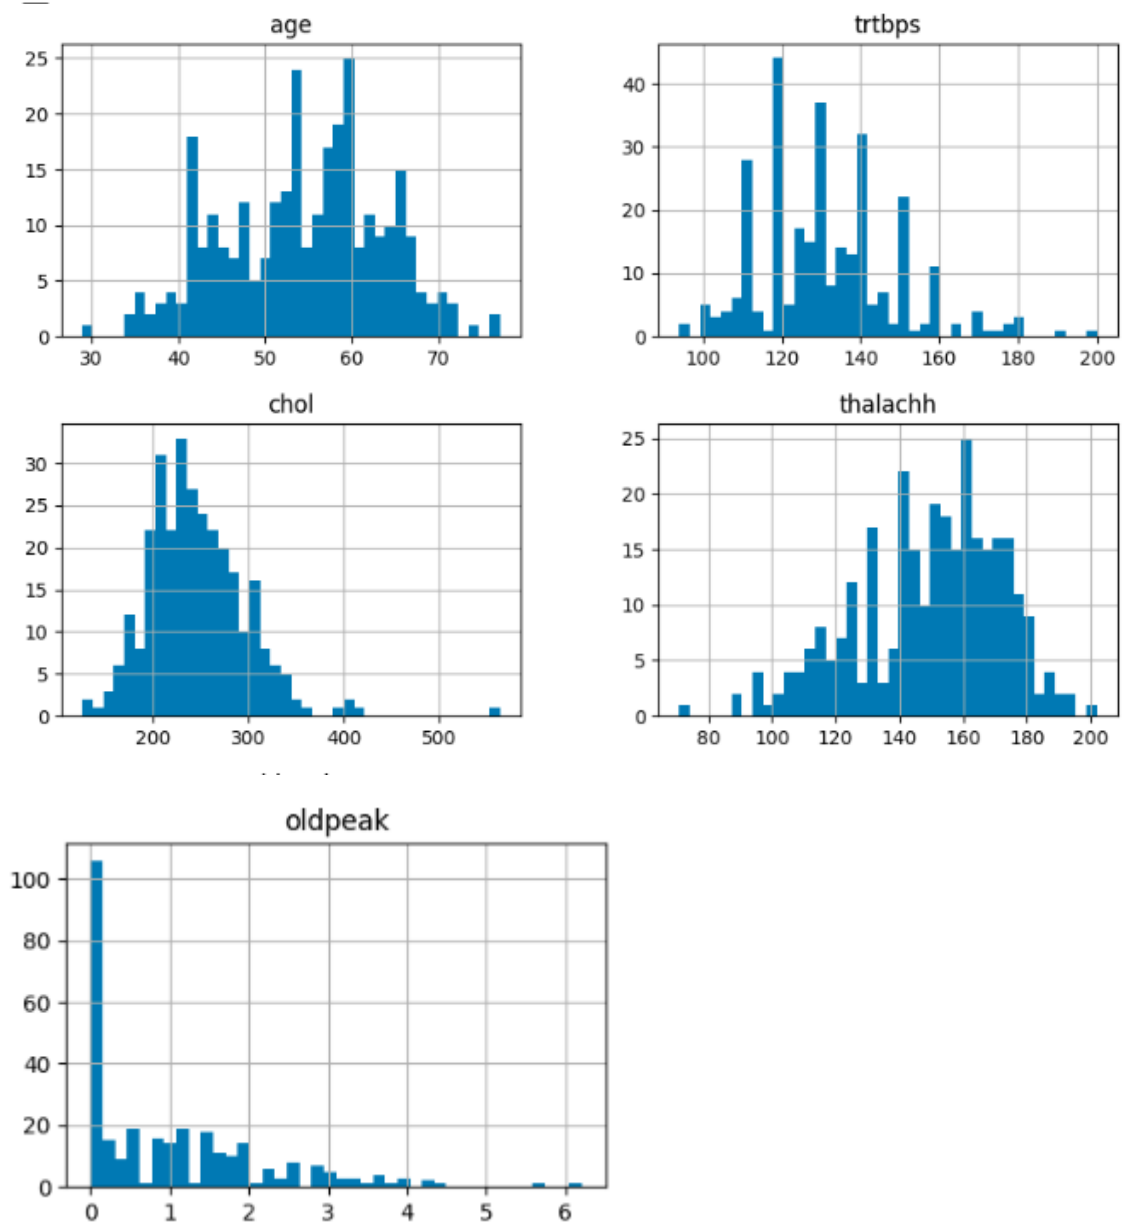
Understand the data better by data visualization.

```
for i in df_const[['sex', 'cp', 'fbs', 'restecg',  
                  'exng', 'slp', 'caa', 'thall', 'output']]:  
    plt.figure(figsize=(3,3))  
    sns.countplot(x=i,data=df_const)  
    plt.show()
```






```
df_const[['age', 'trtbps', 'chol', 'thalachh', 'oldpeak']].hist(bins=40, figsize=(10,10))
plt.show()
```



Correlation heatmap

```
[12] corr = df_const.corr()
plt.figure(figsize=(10,10))
ax = sns.heatmap(corr, vmin=-1, vmax=1, cmap="coolwarm", linewidths=.1, square=True, annot=True, fmt=".2f")
plt.xticks(rotation=0)
plt.show()
```



```
corr['output'].sort_values(ascending=False)

/usr/local/lib/python3.10/dist-packages/ipykernel/ipkerne
and should_run_async(code)
output      1.000000
cp          0.433798
thalachh    0.421741
slp         0.345877
restecg     0.137230
fbs         -0.028046
chol        -0.085239
trtbps      -0.144931
age         -0.225439
sex         -0.280937
thall       -0.344029
caa         -0.391724
oldpeak     -0.430696
exng        -0.436757
Name: output, dtype: float64
```

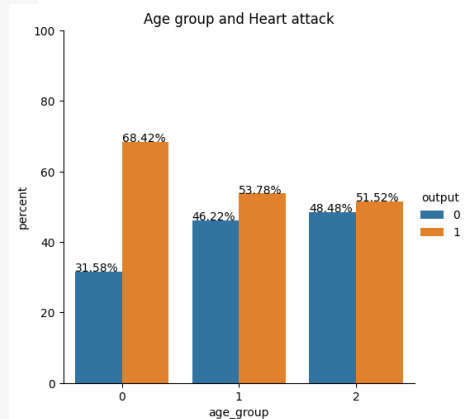
Feature Engineering

The numeric features should be removed and transformed through

➤ Age

```
[78] def encode_age(age):  
    #adults < 40  
    if age <= 40.0:  
        return 0  
    # middle age 41-65  
    elif 40.0 < age <= 65.0:  
        return 1  
    # elderly > 65  
    else:  
        return 2  
  
    # encode the Age to Age_group  
    df['age_group'] = df['age'].apply(encode_age)
```

```
x,y = 'age_group', 'output'  
df1 = df.groupby(x)[y].value_counts(normalize=True)  
df1 = df1.mul(100)  
df1 = df1.rename('percent').reset_index()  
  
g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df1)  
g.ax.set_ylim(0,100)  
plt.title('Age group and Heart attack')  
  
for p in g.ax.patches:  
    txt = str(p.get_height().round(2)) + '%'  
    txt_x = p.get_x()  
    txt_y = p.get_height()  
    g.ax.text(txt_x,txt_y,txt)
```



➤ Maximum Heart rate

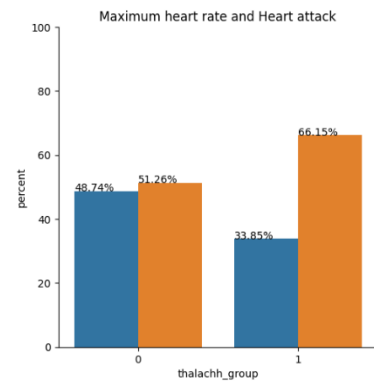
```
[19] def encode_thalachh(age, thalachh):  
    max_hearttrate = 220.0-age  
    # normal  
    if thalachh <= max_hearttrate:  
        return 0  
    # abnormal  
    else:  
        return 1  
  
    # encode the thalachh to thalachh_group  
    df['thalachh_group'] = df.apply(lambda x: encode_thalachh(x['age'], x['thalachh']), axis=1)
```



```
[20] x,y = 'thalachh_group', 'output'
df1 = df.groupby(x)[y].value_counts(normalize=True)
df1 = df1.mul(100)
df1 = df1.rename('percent').reset_index()

g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df1)
g.ax.set_ylim(0,100)
plt.title('Maximum heart rate and Heart attack')

for p in g.ax.patches:
    txt = str(p.get_height().round(2)) + '%'
    txt_x = p.get_x()
    txt_y = p.get_height()
    g.ax.text(txt_x,txt_y,txt)
```



➤ Cholesterol.

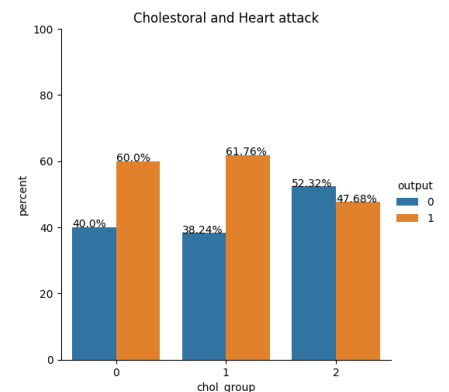
```
[ ] def encode_chol(chol):
    # normal
    if chol < 200.0:
        return 0
    # borderline high
    elif 200.0 <= chol <= 240.0:
        return 1
    # abnormal
    else:
        return 2

    # encode the cholesterol to cholesterol_group
    df['chol_group'] = df['chol'].apply(encode_chol)
```

```
[ ] x,y = 'chol_group', 'output'
df1 = df.groupby(x)[y].value_counts(normalize=True)
df1 = df1.mul(100)
df1 = df1.rename('percent').reset_index()

g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df1)
g.ax.set_ylim(0,100)
plt.title('Cholestoral and Heart attack')

for p in g.ax.patches:
    txt = str(p.get_height().round(2)) + '%'
    txt_x = p.get_x()
    txt_y = p.get_height()
    g.ax.text(txt_x,txt_y,txt)
```



➤ Blood Pressure

```
[2] def encode_trtbps(trtbps):
    # normal
    if trtbps < 120.0:
        return 0
    # pre-hypertension
    elif 120.0 <= trtbps < 140.0:
        return 1
    # hypertension
    else:
        return 2

    # encode the cholesterol to cholesterol_group
    df['trtbps_group'] = df['trtbps'].apply(encode_trtbps)
```

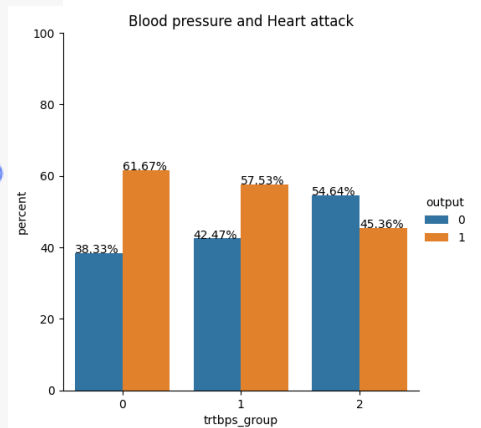
```

x,y = 'trtbps_group', 'output'
df1 = df.groupby(x)[y].value_counts(normalize=True)
df1 = df1.mul(100)
df1 = df1.rename('percent').reset_index()

g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df1)
g.ax.set_ylim(0,100)
plt.title('Blood pressure and Heart attack')

for p in g.ax.patches:
    txt = str(p.get_height().round(2)) + '%'
    txt_x = p.get_x()
    txt_y = p.get_height()
    g.ax.text(txt_x,txt_y,txt)

```



➤ Chest Pain

```

[25] def encode_cp(cp):
    # typical angina
    if cp == 0:
        return 0
    # the other 3
    else:
        return 1

    # encode the cholesterol to cholesterol_group
    df['cp_group'] = df['cp'].apply(encode_cp)

```


```

x,y = 'cp_group', 'output'
df1 = df.groupby(x)[y].value_counts(normalize=True)
df1 = df1.mul(100)
df1 = df1.rename('percent').reset_index()

g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df1)
g.ax.set_ylim(0,100)
plt.title('Chest pain type and Heart attack')

for p in g.ax.patches:
    txt = str(p.get_height().round(2)) + '%'
    txt_x = p.get_x()
    txt_y = p.get_height()
    g.ax.text(txt_x,txt_y,txt)

```



➤ Old Peak

```

def encode_oldpeak(oldpeak):
    # normal
    if oldpeak < 1.0:
        return 0
    # abnormal
    else:
        return 1

    # replace the oldpeak column
    df['oldpeak'] = df['oldpeak'].apply(encode_oldpeak)

```

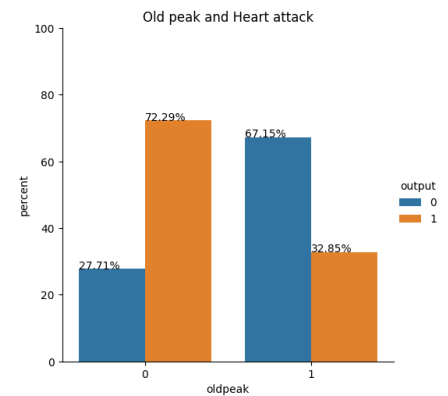
```

x,y = 'oldpeak', 'output'
df1 = df.groupby(x)[y].value_counts(normalize=True)
df1 = df1.mul(100)
df1 = df1.rename('percent').reset_index()

g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df1)
g.ax.set_ylim(0,100)
plt.title('Old peak and Heart attack')

for p in g.ax.patches:
    txt = str(p.get_height().round(2)) + '%'
    txt_x = p.get_x()
    txt_y = p.get_height()
    g.ax.text(txt_x,txt_y,txt)

```



+ Drop unwanted column

```

df.drop(columns=['age', 'chol', 'thalachh', 'trtbps'], inplace=True)

[30] df.head()

```

/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell` and should_run_async(code)

	sex	cp	fbs	restecg	exng	oldpeak	slp	caa	thall	output	age_group	thalachh_group	chol_group	trtbps_group	cp_group
0	1	3	1	0	0	1	0	0	1	1	1	0	1	2	1
1	1	2	0	1	0	1	0	0	2	1	0	1	2	1	1
2	0	1	0	0	0	1	2	0	2	1	1	0	1	1	1
3	1	1	0	1	0	0	2	0	2	1	1	1	1	1	1
4	0	0	0	1	1	0	2	0	2	1	1	0	2	1	0

+ Feature selection by Association Rule Mining

➤ Data Preparation

```

df_nohead=df.copy(deep=True)

# add column as prefix for each value to indentify them in rules
for x in df_nohead.columns.values.tolist():
    df_nohead[x] = x + '_' + df_nohead[x].astype(str)

df_nohead.columns = range(df_nohead.shape[1])
df_nohead

```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
0	sex_1	cp_3	fbs_1	restecg_0	exng_0	oldpeak_1	slp_0	caa_0	thall_1	output_1	age_group_1	thalachh_group_0	chol_group_1	trtbps_group_2	cp_group_1
1	sex_1	cp_2	fbs_0	restecg_1	exng_0	oldpeak_1	slp_0	caa_0	thall_2	output_1	age_group_0	thalachh_group_1	chol_group_2	trtbps_group_1	cp_group_1
2	sex_0	cp_1	fbs_0	restecg_0	exng_0	oldpeak_1	slp_2	caa_0	thall_2	output_1	age_group_1	thalachh_group_0	chol_group_1	trtbps_group_1	cp_group_1
3	sex_1	cp_1	fbs_0	restecg_1	exng_0	oldpeak_0	slp_2	caa_0	thall_2	output_1	age_group_1	thalachh_group_1	chol_group_1	trtbps_group_1	cp_group_1
4	sex_0	cp_0	fbs_0	restecg_1	exng_1	oldpeak_0	slp_2	caa_0	thall_2	output_1	age_group_1	thalachh_group_0	chol_group_2	trtbps_group_1	cp_group_0
...
298	sex_0	cp_0	fbs_0	restecg_1	exng_1	oldpeak_0	slp_1	caa_0	thall_3	output_0	age_group_1	thalachh_group_0	chol_group_2	trtbps_group_2	cp_group_0
299	sex_1	cp_3	fbs_0	restecg_1	exng_0	oldpeak_1	slp_1	caa_0	thall_3	output_0	age_group_1	thalachh_group_0	chol_group_2	trtbps_group_0	cp_group_1
300	sex_1	cp_0	fbs_1	restecg_1	exng_0	oldpeak_1	slp_1	caa_2	thall_3	output_0	age_group_2	thalachh_group_0	chol_group_0	trtbps_group_2	cp_group_0
301	sex_1	cp_0	fbs_0	restecg_1	exng_1	oldpeak_1	slp_1	caa_1	thall_3	output_0	age_group_1	thalachh_group_0	chol_group_0	trtbps_group_1	cp_group_0
302	sex_0	cp_1	fbs_0	restecg_0	exng_0	oldpeak_0	slp_1	caa_1	thall_2	output_0	age_group_1	thalachh_group_1	chol_group_1	trtbps_group_1	cp_group_1

```

te = TransactionEncoder()
transactions = df_nohead.values.tolist()
te_ary = te.fit_transform(transactions)
result_df = pd.DataFrame(te_ary, columns=te.columns_)
result_df

```

	age_group_0	age_group_1	age_group_2	caa_0	caa_1	caa_2	caa_3	caa_4	chol_group_0	chol_group_1	...	slp_2	thalachh_group_0	thalachh_group_1	thall_0	thall_1	thall_2	thall_3	trtbps_group_0	trtbps_group_1	trtbps_group_2
0	False	True	False	True	False	False	False	False	False	True	...	False	True	False	False	True	False	False	False	False	True
1	True	False	False	True	False	False	False	False	False	False	...	False	False	True	False	False	True	False	False	True	False
2	False	True	False	True	False	False	False	False	False	True	...	True	True	False	False	False	True	False	False	True	False
3	False	True	False	True	False	False	False	False	False	True	...	True	False	True	False	False	True	False	False	True	False
4	False	True	False	True	False	False	False	False	False	False	...	True	True	False	False	False	True	False	False	True	False
...
298	False	True	False	True	False	False	False	False	False	False	...	False	True	False	False	False	False	True	False	False	True
299	False	True	False	True	False	False	False	False	False	False	...	False	True	False	False	False	False	True	True	False	False
300	False	False	True	False	False	True	False	False	True	False	...	False	True	False	False	False	False	True	False	False	True
301	False	True	False	False	True	False	False	False	True	False	...	False	True	False	False	False	False	True	False	True	False
302	False	True	False	False	True	False	False	False	False	True	...	False	False	True	False	False	False	True	False	True	False

➤ Apriori Algorithm

```

freq_items = apriori(result_df, min_support=0.3, use_colnames=True)
freq_items.head(10)

```

```

/usr/local/lib/python3.10/dist-packages/ipynbkernel/ipkernel.py:283: I
and should_run_async(code)

```

	support	itemsets
0	0.828383	(age_group_1)
1	0.577558	(caa_0)
2	0.336634	(chol_group_1)
3	0.498350	(chol_group_2)
4	0.471947	(cp_0)
5	0.471947	(cp_group_0)
6	0.528053	(cp_group_1)
7	0.673267	(exng_0)
8	0.326733	(exng_1)
9	0.851485	(fbs_0)

```
rules = association_rules(freq_items, metric="support", min_threshold=0.3)
rules.sort_values("support", ascending=False).head(5)
```

```
/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell` automatically and should_run_async(code)
```

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction	zhangs_metric
13	(age_group_1)	(fbs_0)	0.828383	0.851485	0.699670	0.844622	0.991939	-0.005686	0.955826	-0.045210
12	(fbs_0)	(age_group_1)	0.851485	0.828383	0.699670	0.821705	0.991939	-0.005686	0.962548	-0.051878
140	(fbs_0)	(thalachh_group_0)	0.851485	0.785479	0.676568	0.794574	1.011579	0.007744	1.044274	0.077073
141	(thalachh_group_0)	(fbs_0)	0.785479	0.851485	0.676568	0.861345	1.011579	0.007744	1.071107	0.053358
32	(thalachh_group_0)	(age_group_1)	0.785479	0.828383	0.650165	0.827731	0.999213	-0.000512	0.996217	-0.003657

➤ Data Filtering and visualization

```
[36] rules['consequents_len'] = rules['consequents'].apply(lambda x: len(x))
rules
```

```
[37] sorted_rules = rules[(rules['consequents_len'] == 1)]
sorted_rules = sorted_rules.sort_values("lift", ascending=False)
```

```
[38] sorted_rules.drop_duplicates(subset=['antecedents', 'consequents'], inplace=True)
```

```
[39] sorted_rules[sorted_rules['consequents'].astype(str).str.contains("output")].head(10)
```

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction	zhangs_metric	consequents_len
634	(caa_0, thall_2)	(output_1)	0.376238	0.544554	0.336634	0.894737	1.643062	0.131752	4.326733	0.627451	1
753	(cp_group_1, thall_2)	(output_1)	0.376238	0.544554	0.333333	0.885965	1.626954	0.128451	3.993907	0.617790	1
1186	(fbs_0, caa_0, thall_2)	(output_1)	0.343234	0.544554	0.303630	0.884615	1.624476	0.116721	3.947195	0.585318	1
1242	(cp_group_1, exng_0, thall_2)	(output_1)	0.343234	0.544554	0.300330	0.875000	1.606818	0.113420	3.643564	0.575018	1
555	(cp_group_1, caa_0)	(output_1)	0.363036	0.544554	0.316832	0.872727	1.602645	0.119139	3.578501	0.590350	1
80	(cp_group_0)	(output_0)	0.471947	0.455446	0.343234	0.727273	1.596838	0.128288	1.996700	0.707813	1
72	(cp_0)	(output_0)	0.471947	0.455446	0.343234	0.727273	1.596838	0.128288	1.996700	0.707813	1
650	(cp_0, cp_group_0)	(output_0)	0.471947	0.455446	0.343234	0.727273	1.596838	0.128288	1.996700	0.707813	1
991	(exng_0, age_group_1, caa_0)	(output_1)	0.359736	0.544554	0.310231	0.862385	1.583653	0.114335	3.309571	0.575620	1
574	(exng_0, caa_0)	(output_1)	0.432343	0.544554	0.369637	0.854962	1.570021	0.134203	3.140177	0.639587	1

```
sorted_rules[(sorted_rules['antecedents'].astype(str).str.contains("thalachh")
| sorted_rules['antecedents'].astype(str).str.contains("slp")
| sorted_rules['antecedents'].astype(str).str.contains("restecg"))
& sorted_rules['consequents'].astype(str).str.contains("output")].head(10)
```

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction	zhangs_metric	consequents_len
820	(exng_0, slp_2)	(output_1)	0.382838	0.544554	0.310231	0.810345	1.488088	0.101755	2.401440	0.531460	1
969	(thalachh_group_0, thall_2)	(output_1)	0.392739	0.544554	0.306931	0.781513	1.435141	0.093063	2.084539	0.499299	1
169	(slp_1)	(output_0)	0.462046	0.455446	0.300330	0.650000	1.427174	0.089893	1.555870	0.556395	1
177	(slp_2)	(output_1)	0.468647	0.544554	0.353135	0.753521	1.383739	0.097932	1.847808	0.521913	1
903	(fbs_0, slp_2)	(output_1)	0.402640	0.544554	0.300330	0.745902	1.369747	0.081070	1.792399	0.451885	1
963	(thalachh_group_0, sex_1)	(output_0)	0.531353	0.455446	0.320132	0.602484	1.322846	0.078130	1.369895	0.520764	1
1255	(fbs_0, thalachh_group_0, exng_0)	(output_1)	0.429043	0.544554	0.300330	0.700000	1.285455	0.066693	1.518152	0.388935	1
627	(thalachh_group_0, caa_0)	(output_1)	0.448845	0.544554	0.313531	0.698529	1.282754	0.069111	1.510746	0.399937	1
825	(thalachh_group_0, exng_0)	(output_1)	0.488449	0.544554	0.339934	0.695946	1.278010	0.073947	1.497910	0.425243	1
173	(restecg_1)	(output_1)	0.501650	0.544554	0.316832	0.631579	1.159809	0.043656	1.236209	0.276490	1

The best features selected by ARM are: caa, thall, cp_group, fbs, exng, cp, age_group

Modelling

```
validation_results = pd.DataFrame(columns=["model", "accuracy"])
```

Data Preparation

```
[43] df_final_arm = df[['caa', 'thall', 'cp_group', 'fbs', 'exng', 'cp', 'age_group', 'output']].copy(deep=True)
      df_final_corr = df[['cp_group', 'cp', 'slp', 'restecg', 'thalachh_group', 'output']].copy(deep=True)

      # final dataset of original/ unprocessed data
      df_final_original = df_const.copy(deep=True)
```

```
df_final = df_final_arm.copy(deep=True)
```

```
[47] X = df_final.drop(columns=['output'])
      y = df['output']

      X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=0.25, random_state=0)
      X_train
```

	caa	thall	cp_group	fbs	exng	cp	age_group
173	2	3	1	0	0	2	1
261	1	2	0	0	0	0	1
37	0	3	1	0	0	2	1
101	0	3	1	0	0	3	1
166	2	3	0	0	1	0	2
...
251	4	3	0	1	1	0	1
192	1	3	0	0	0	0	1
117	0	3	1	0	0	3	1
47	0	2	1	0	0	2	1
172	0	2	1	0	0	1	1

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 227 entries, 173 to 172
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   caa          227 non-null   int64  
1   thall        227 non-null   int64  
2   cp_group     227 non-null   int64  
3   fbs          227 non-null   int64  
4   exng         227 non-null   int64  
5   cp           227 non-null   int64  
6   age_group    227 non-null   int64  
dtypes: int64(7)
```

Classification

```
[49] def create_model(units=128, activation='relu', learning_rate=0.001, optimizer='adam'):
    model = Sequential()
    model.add(Dense(units, input_dim=X_train.shape[1], activation=activation))
    model.add(Dense(1, activation='sigmoid')) # Assuming binary classification
    model.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])
    return model
```

```
estimator = {
    'RandomForest': RandomForestClassifier(random_state=0),
    'DecisionTree': DecisionTreeClassifier(),
    'Knn': KNeighborsClassifier(),
    'Keras': KerasClassifier(model=create_model, verbose=0, random_state=0),
    'LogisticRegression': LogisticRegression(max_iter=100000, random_state=0),
}

param_grid = {
    'RandomForest': {
        'n_estimators': range(10, 110, 10),
        'max_depth': range(3, 22, 2),
    },
    'DecisionTree': {
        'max_depth': range(3, 22, 2),
    },
    'Knn': {
        'n_neighbors': range(3, 110, 3),
        'weights': ['uniform', 'distance'],
    },
    'Keras': {
        'optimizer': ['adam', 'rmsprop'],
        'batch_size': [32, 64],
        'epochs': [10, 20, 30, 40, 50],
    },
    'LogisticRegression': {
        'solver': ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'],
    },
}
```

```
fit_models = {}
# loop all the classifiers
for algo, classifier in estimator.items():
    print(f'Training the {algo} model...')
    # create grid search CV
    clf = GridSearchCV(estimator=classifier, param_grid=param_grid[algo], n_jobs=-1, cv=None)

    # train the model
    t0 = time.time()
    clf.fit(X_train, y_train)

    # store the trained result
    fit_models[algo] = clf

    print(f'The {algo} model is trained. Time taken = {time.time()-t0}')
```

```
Training the RandomForest model...
/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: `should_run_async` will not call `transform_cell` automa
and should_run_async(code)
The RandomForest model is trained. Time taken = 48.503103494644165
Training the DecisionTree model...
The DecisionTree model is trained. Time taken = 0.22644901275634766
Training the Knn model...
The Knn model is trained. Time taken = 1.9948561191558838
Training the Keras model...
/usr/local/lib/python3.10/dist-packages/joblib/externals/loky/process_executor.py:752: UserWarning: A worker stopped while some jobs were given
warnings.warn(
The Keras model is trained. Time taken = 117.9610641002655
Training the LogisticRegression model...
The LogisticRegression model is trained. Time taken = 0.2695605754852295
```

```

▶ for algo, classifier in fit_models.items():
    y_pred = classifier.predict(X_valid)
    accuracy = accuracy_score(y_valid, y_pred)
    print(f'Model: {algo}')
    print(f'Accuracy score: {accuracy}\n')
    validation_results.loc[len(validation_results.index)] = [algo, accuracy]

```

➞ /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: DeprecationWarning: and should_run_async(code)

Model: RandomForest
Accuracy score: 0.8289473684210527

Model: DecisionTree
Accuracy score: 0.8289473684210527

Model: Knn
Accuracy score: 0.8552631578947368

Model: Keras
Accuracy score: 0.7763157894736842

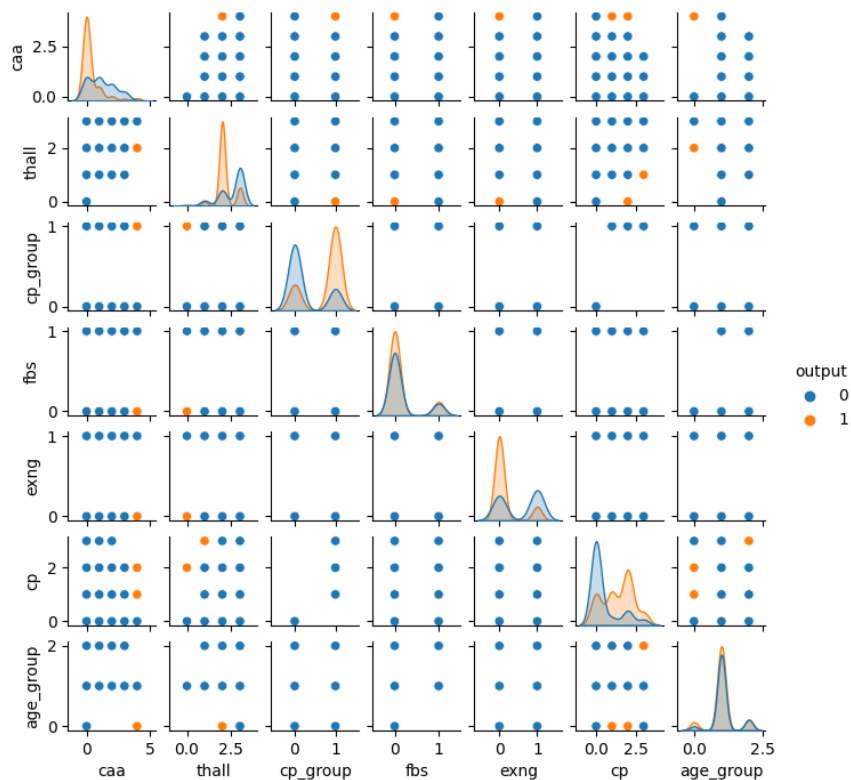
Model: LogisticRegression
Accuracy score: 0.8157894736842105

✚ Clustering

```

✓ 18s [54] X_cluster=X_train.copy(deep=True)
      sns.pairplot(df_final, hue='output', height=1)
      plt.show()

```

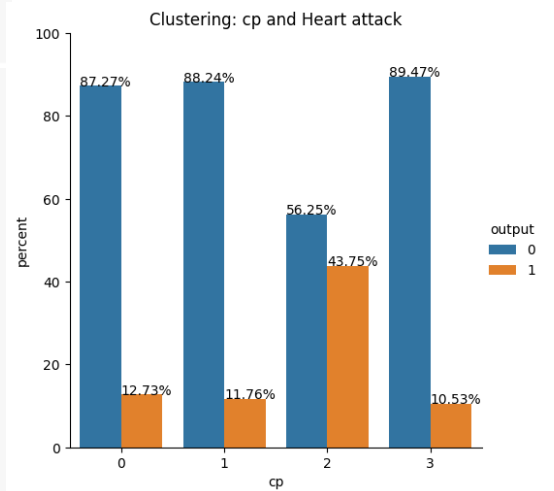



```
[63] km=KModes(n_clusters=2)
      cluster_labels=km.fit_predict(X_cluster)
      X_cluster['output'] = cluster_labels
```

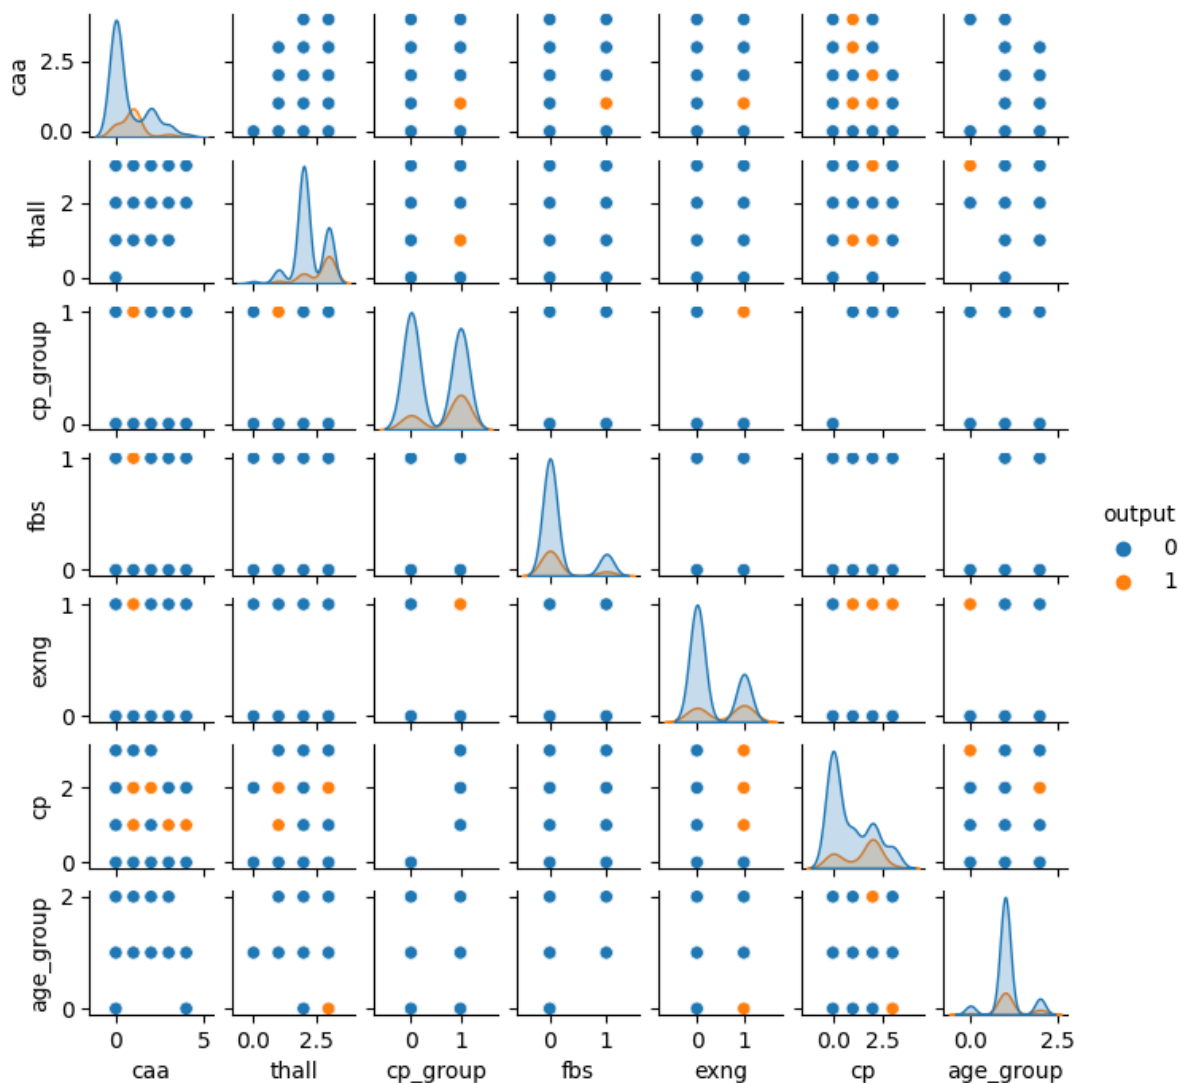
```
▶ x,y = 'cp', 'output'
  df9 = X_cluster.groupby(x)[y].value_counts(normalize=True)
  df9 = df9.mul(100)
  df9 = df9.rename('percent').reset_index()

  g = sns.catplot(x=x,y='percent',hue=y,kind='bar',data=df9)
  g.ax.set_ylim(0,100)
  plt.title('Clustering: cp and Heart attack')

  for p in g.ax.patches:
      txt = str(p.get_height().round(2)) + '%'
      txt_x = p.get_x()
      txt_y = p.get_height()
      g.ax.text(txt_x,txt_y,txt)
```



```
[65] sns.pairplot(X_cluster, hue='output', height=1)
      plt.show()
```



Results Visualization

```
[5] validation_results
```

```
/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283:  
and should_run_async(code)
```

	model	accuracy
0	RandomForest	0.828947
1	DecisionTree	0.828947
2	Knn	0.855263
3	Keras	0.776316
4	LogisticRegression	0.815789

➤ Sorting results by Accuracy

```
[69] validation_results.sort_values(by="accuracy", ascending=False)
```

```
/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283:  
and should_run_async(code)
```

	model	accuracy
2	Knn	0.855263
0	RandomForest	0.828947
1	DecisionTree	0.828947
4	LogisticRegression	0.815789
3	Keras	0.776316